

AMAN NEGI

amanpy45@gmail.com • +91 9045473895 • Uttarakhand, India

[GitHub](#) | [LinkedIn](#) | [LeetCode](#)

Computer Science student with hands-on experience across the full ML lifecycle — from data pipelines and statistical modeling to deep learning and deployment. Currently applying these skills as a Data Science Intern, with deployed deep learning projects in computer vision and NLP.

EDUCATION

DIT University — B.Tech in Computer Science and Engineering, CGPA: 7.4

Aug 2023 – May 2027

Class XII (Secondary) — CGPA: 7.89 | Class X (Senior Secondary) — CGPA: 8.86

EXPERIENCE

Data Science Intern — Celebal Technologies

Jun 2026 – Present · Remote

- Built an end-to-end regression and time-series forecasting pipeline on real-world Tesla production/delivery data using Linear Regression, Random Forest, and ARIMA; tuned models with GridSearchCV and evaluated using MAE, RMSE, and R-squared.
- Performed country segmentation on socio-economic indicators combining unsupervised (K-Means, DBSCAN) and supervised (Decision Tree, Random Forest, XGBoost) methods; validated clusters via Elbow Method, Silhouette Score, and PCA.
- Designed and trained ANN and CNN architectures on the CIFAR-10 dataset, benchmarking accuracy/loss to evaluate convolutional vs. fully-connected approaches for image classification.
- Applied core statistical foundations (hypothesis testing, CLT, Bayes' theorem, PSI) and model evaluation metrics to real-world data analysis tasks.

PROJECTS

Deepfake Image Detector

TensorFlow, MobileNetV2, Streamlit

[GitHub](#) | [Live Demo](#)

- Built a binary real-vs-deepfake image classifier using transfer learning with a frozen, ImageNet-pretrained MobileNetV2 backbone as a feature extractor.
- Designed a lightweight custom classification head (GlobalAveragePooling2D → Dense(32, ReLU) → Dense(1, sigmoid)), keeping the pipeline fully CPU-trainable.
- Deployed the trained model as a live, interactive Streamlit web app supporting real-time image upload with real/fake prediction and confidence score.

Sentiment Analysis — Bi-LSTM & Transformer Benchmark

TensorFlow, NLP, Bi-LSTM

[GitHub](#) | [Live Demo](#)

- Built a Bidirectional LSTM model to classify ~61,000+ tweets into Positive, Negative, or Neutral sentiment using the Twitter Entity Sentiment dataset.
- Achieved 98% accuracy and a 0.98 macro F1-score on the validation set (828 samples), with balanced precision/recall across all three classes.
- Benchmarked the trained model against a pretrained RoBERTa transformer (zero-shot) to compare domain-specific vs. general-purpose language understanding on negation and contextual sentiment.
- Deployed the model as a live Streamlit app for real-time sentiment prediction on user-entered text.

TECHNICAL SKILLS

Languages: Python, Java, SQL

ML / DL: TensorFlow, Keras, PyTorch, Scikit-learn, NumPy, Pandas

Tools: Git, GitHub, Streamlit, Jupyter, VS Code, MySQL, Linux

Coding Practice: 255+ problems solved on LeetCode (Python, Java, MySQL) · 100-day solving streak

CERTIFICATIONS & ACHIEVEMENTS

- **Python Developer** — HackerRank Certified
- **Machine Learning & Deep Learning** — Professional Specializations via Coursera
- **Database Management** — SQL Certification (Coursera)
- **Dataset Contributor** — Published ML-ready datasets on Kaggle for public research

LEADERSHIP & INITIATIVE

- Mentored students in machine learning, neural networks, and data structures through peer learning sessions.
- Explored AI/analytics applications via self-driven experimentation and open-source collaboration.